Binary Tree Classifier from Scratch for Mushroom Classification

Final Project in the Subject Machine Learning

Julia Maria Wdowinska

Data Science for Economics
I year
Master's Degree
Matriculation Number: 43288A

January 26, 2025

Contents

1 Introduction

The classification of mushrooms as edible or poisonous based on their features is an interesting task. Although such a model might not be sufficiently precise to definitively instruct a person whether to pick or avoid a particular mushroom, it can serve as a valuable aid or provide confirmation that a person's intuition about a specific mushroom is correct.

This project implements a custom-built decision tree classifier designed to solve this (and other binary) classification problem. The goal of the implementation was to ensure simplicity, efficiency, and flexibility, offering various splitting and stopping criteria, options for creating splits based on numerical and categorical variables, and the capability to handle missing values.

2 Tree Classifier Implementation

Tree predictors are fundamental tools in machine learning, widely applied to classification and regression tasks. They represent a hierarchy of decision rules, where data points are recursively split into subsets based on feature values. The main advantage of tree predictors is their ability to handle both numerical and categorical features. Tree predictors are also straightforward, making them a popular choice when interpretability is a priority.

In this study, a complete binary tree classifier - where each internal node has exactly two children - has been implemented in Python. A detailed description of the classes and methods created is provided below.

2.1 TreeNode Class

The `TreeNode` class represents a single node in a binary tree classifier. Each node can either be an internal node or a leaf node. Internal nodes split the data based on a specific feature and threshold, while leaf nodes store the predicted label. The attributes and methods of this class are designed to support the recursive structure of the tree classifier.

Attributes:

- `feature_index` (int or None): The feature index used for splitting the data at this node.

- `threshold_value` (float, str or None): The threshold value used for splitting the data at this node.

- `left_child` (TreeNode or None): The left child node.

- `right_child` (TreeNode or None): The right child node.

- `left_ratio` (float or None): The ratio of samples that go to the left child. This value is particularly useful for handling missing values and calculating probabilities.

- `leaf_value` (int or None): The predicted label associated with the leaf node.

If the node is a leaf node, then `feature_index`, `threshold_value`, `left_child`, `right_child`, and `left_ratio` are None, while `leaf_value` is an integer. If the node is not a leaf node, then only `leaf_value` is None.

Methods:

- `is_leaf()`: This method checks whether the current node is a leaf node. It returns `True` if the node has a `leaf_value`, and `False` otherwise.

2.2 DecisionTreeClassifier Class

The `DecisionTreeClassifier` class implements a decision tree for binary classification. It recursively splits the data based on specific features and thresholds, creating a tree structure that can be used for predicting labels. The attributes and methods of this class are designed to support model training, hyperparameter tuning, and prediction.

Attributes:

- `min_samples_split` (int): The minimum number of samples required to split a node.

- `max_depth` (int or None): The maximum depth of the tree. If set to `None`, the tree expands until all nodes are pure, contain fewer than `min_samples_split` samples, or further splitting results in an information gain below `min_information_gain`.

- `n_features` (int, float, or str): The number of features to consider when identifying the best split. This can be specified as an integer, a float, or one of the following strings: 'sqrt' or 'log2'.

- `criterion` (str): The function to measure the quality of a split. Options are: 'gini', 'scaled_entropy', and 'square_root'.

- `min_information_gain` (float): The minimum information gain required to perform a split.

- `n_quantiles` (int or None): The number of quantiles to consider when determining the best threshold for continuous features. If set to `None`, the algorithm uses midpoints of unique values.

- `isolate_one` (bool): Whether to isolate a single value for categorical features, creating a one-vs-rest split.

- `root` (TreeNode or None): The root node of the decision tree.

- `depth` (int): The final depth of the tree after it has been built.

min_samples_split Parameter:

The `min_samples_split` parameter controls the minimum number of samples a node must contain to be eligible for splitting. If a node has fewer than `min_samples_split` samples, it becomes a leaf node, and no further splits are attempted.

A higher value for `min_samples_split` reduces the depth of the tree, making it less prone to capturing noise in the data. Conversely, a lower value allows the tree to grow deeper and potentially capture finer details, which can be beneficial for highly complex datasets but may increase the risk of overfitting. The default value is 2.

n_features Parameter:

The `n_features` parameter specifies the number of features to consider when identifying the optimal split. The default value is `None`, which results in all features in the dataset being considered. Otherwise:

- If `n_features` is an integer, this specifies the number of features to consider. If the value exceeds the total number of features in the dataset, all features are considered instead.

- If `n_features` is a float, it represents a fraction of the total number of features. The number of features to consider is calculated by multiplying this fraction by the total number of features, and truncating the decimal part to obtain an integer. At least one feature is considered.

- If `n_features` is a string, it can be either 'sqrt' or 'log2', and the number of features is calculated as follows:

    - 'sqrt': Sets the number of features to the square root of the total number of features, truncating the decimal part to obtain an integer. At least one feature is considered.

    - 'log2': Sets the number of features to the base-2 logarithm of the total number of features, truncating the decimal part to obtain an integer. At least one feature is considered.

This parameter allows the decision tree model to use a subset of features, which can help improve the model's efficiency and performance, particularly when working with high-dimensional datasets.

criterion Parameter:

The `criterion` parameter specifies the function used to measure the quality of a split in the decision tree. The available options are:

- '`gini`' (the default): The Gini impurity is used, which is computed as:

$$\text{Gini} = 2 \cdot p_0 \cdot (1 - p_0)$$

  where p_0 is the probability of class '0' within the node.

- '`scaled_entropy`': The scaled entropy is used. The entropy is scaled by halving the probabilities before applying the standard entropy formula:

$$\text{Scaled Entropy} = -\sum_i \frac{p_i}{2} \cdot \log_2(p_i + \epsilon)$$

  where p_i is the probability of class i, and ϵ is a small constant to avoid taking the logarithm of zero.

- '`square_root`': The "square root" impurity is used, which is calculated as:

$$\text{Square Root Impurity} = \sqrt{p_0 \cdot (1 - p_0)}$$

  where p_0 is the probability of class '0' within the node.

min_information_gain Parameter:

The `min_information_gain` parameter specifies the minimum amount of information gain required to perform a split. Information gain measures the reduction in impurity after a split. It is computed as follows:

$$\text{Information Gain} = \text{Impurity Before Split} - \text{Weighted Impurity After Split}$$

where the impurity is calculated using the selected `criterion`, such as Gini impurity, scaled entropy, or square root impurity. The weighted impurity after the split is calculated as:

$$\text{Weighted Impurity After Split} = \frac{L}{n} \cdot \text{Impurity of Left Child} + \frac{R}{n} \cdot \text{Impurity of Right Child}$$

where:

- L and R are the number of samples in the left and right child nodes, respectively.

- n is the total number of samples in the parent node.

The `min_information_gain` parameter accepts a float value that sets the threshold for the minimum information gain. If the calculated information gain from a potential split is less than this threshold, the split is not performed, and the node becomes a leaf node. The default value is 0.0.

n_quantiles Parameter:

The `n_quantiles` parameter determines how candidate thresholds are chosen when splitting based on numerical features. If set to `None` (the default), all midpoints between unique values are considered. Otherwise:

- If `n_quantiles` is an integer, the values are divided into that many quantiles, and the candidate thresholds are the boundaries between these quantiles.

While lower values of `n_quantiles` reduce the number of candidate thresholds, speeding up computation but potentially leading to suboptimal splits, higher values or setting it to `None` (to consider all midpoints) increase the search granularity, increasing the probability of finding an optimal split, but at the cost of additional computation time.

isolate_one Parameter:

The `isolate_one` parameter determines how splits are made when splitting based on categorical features. If set to `False` (the default), all data points with a feature value lower or equal (i.e., lower or equal alphabetically) to the threshold are assigned to the left child, while all other data points are assigned to the right child. Otherwise:

- If `isolate_one` is set to `True`, the algorithm creates a one-vs-rest split, where all data points with a feature value equal to the threshold go to the left child, while all other data points go to the right child.

This parameter affects the granularity of splits for categorical features. Setting `isolate_one` to `True` results in more precise splits, capturing finer patterns in the data but potentially increasing the risk of overfitting. In contrast, setting it to `False` produces broader, more generalized splits, improving computational efficiency and helping reduce overfitting.

Private Methods:

- `_build_tree()`: Recursively builds the tree by splitting the data based on the best feature and threshold. It stops if any stopping condition is met, e.g., `max_depth`.

- `_get_most_common_label()`: This method finds and returns the most common label in a given array.

- `_find_best_split()`: Finds the best feature and threshold for splitting the data.

- `_calculate_information_gain()`: Computes the information gain from a potential split based on a selected criterion.

- `_split()`: Splits the data based on the selected feature and threshold.

- `_gini_impurity()`: Computes the Gini impurity for the given labels.

- `_scaled_entropy()`: Computes the scaled entropy for the given labels.

- `_square_root_impurity()`: Computes the "square root" impurity for the given labels.

- `_traverse_tree()`: Traverses the tree for a single input sample and returns the predicted label.

Public Methods:

- `fit()`: Initializes the root node and builds the tree using the `_build_tree()` method.

- `predict()`: Predicts the labels for the given input samples by traversing the tree for each sample using the `_traverse_tree()` method.

_build_tree() Method:

The `_build_tree()` method constructs a decision tree by starting at the root node and progressing recursively to the leaf nodes. At each node, it selects a random subset of features, as specified by the `n_features` parameter, and uses the `_find_best_split()` method to determine the optimal feature and threshold for splitting the data. Once the optimal split is found, the `_split()` method is called to partition the data accordingly. This process is then repeated recursively on the resulting subsets to continue building the tree.

The recursion halts when a stopping condition is met, such as reaching the `max_depth`, achieving pure nodes, having fewer than `min_samples_split` samples per node, or when subsequent splits yield information gains lower than `min_information_gain`. Upon termination, the method assigns the most frequent label among the samples at a node as the predicted label for that node.

_find_best_split() Method:

The `_find_best_split()` method identifies the optimal split for a given node in the tree being built. It iterates through all features selected by the `_build_tree()` method and evaluates all candidate thresholds by calling the `_calculate_information_gain()` method to determine the feature-threshold combination that maximizes information gain. The creation of candidate thresholds differs between numerical and categorical features:

- For numerical features, potential thresholds are determined based on the `n_quantiles` parameter (as described above).

- For categorical features, candidate thresholds consist of all unique values in the feature.

Any missing values are ignored when determining the candidate thresholds.

_split() Method:

The `_split()` method partitions the data based on a specified feature and threshold. The partitioning strategy differs for numerical and categorical features:

- For numerical features, data points with a feature value lower than or equal to the threshold are assigned to the left child, while all other data points are assigned to the right child.

- For categorical features, the partitioning depends on the `isolate_one` parameter (as described above).

After the split, any data points with a missing feature value are randomly distributed between the left and right child. The probability of being assigned to each child is proportional to the number of data points assigned to that child during the split.

_traverse_tree() Method:

The `_traverse_tree()` method traverses the decision tree to predict the label for a single instance. Starting at the root node, it follows the tree's decision rules until it reaches a leaf node.

If a feature value at the current node is missing, the method randomly decides whether to move to the left or right child, based on the ratio of data points assigned to each child.

2.3   Training and Validation Procedures

Training and testing a model on two separate datasets is a fundamental practice in machine learning. The goal is to develop a model that generalizes well to unseen data, and not just memorizes the data seen during training (a phenomenon known as overfitting).

In this study, a set of custom functions has been implemented to facilitate robust model training, validation, and hyperparameter tuning.

Functions:

- `k_fold_partition()`: This function partitions the data into k folds for k-fold cross-validation. The number of folds is specified by the user (the default is 5).

- `k_fold_cv_estimate()`: This function uses the `k_fold_partition()` function to create k folds and trains the model k times, each time using $k-1$ folds for training and the remaining fold for testing. The average test error across all iterations (i.e., the cross-validation estimate) is returned.

- `hyperparameter_tuning()`: This function iterates through all parameter combinations generated by the `_parameter_combinations()` function, or a random subset if specified, and uses the `k_fold_cv_estimate()` function to compute the cross-validation estimate for a model with these parameters. The lowest cross-validation estimate and the corresponding parameters are then returned.

- `k_fold_nested_cv()`: This function implements nested cross-validation. The `k_fold_partition()` function is used to create k folds and the model is trained k times. Each time, the `hyperparameter_tuning()` function is invoked on $k-1$ folds and the model with the best parameters is tested on the remaining fold. A list of k best parameters is then returned, along with a list of all metrics (test error, precision, etc.).

- `_parameter_combinations()`: This helper function generates all possible combinations of hyperparameters from a specified grid.

- `accuracy_metric()`: This function computes the accuracy of the model's predictions.

- `precision_metric()`: This function computes the precision of the model's predictions.

- `recall_metric()`: This function computes the recall of the model's predictions.

- `f1_metric()`: This function computes the F1 score, the harmonic mean of precision and recall.

3   Dataset Description and Preprocessing

The dataset used in this project is a simulated version inspired by the Mushroom Data Set. It contains 61,069 hypothetical mushrooms, each described by 20 features and classified as either definitely edible or definitely poisonous/of unknown edibility. The class distribution is balanced, with 27,181 mushrooms (44.51%) classified as edible and 33,888 (55.49%) as poisonous, ensuring no class is overrepresented in subsequent analyses.

During preprocessing, 146 duplicate rows (0.24%) were identified and removed to avoid undue weighting of observations. This step reduced the dataset to 60,923 unique mushrooms, consisting of 27,181 (44.62%) edible and 33,742 (55.38%) poisonous instances. Additionally, nine variables were found to contain missing values, as summarized in Table 1.

Upon examining the possible values for each variable, some unexpected values not defined by the dataset's author were observed, such as "d" for cap-surface and "f" for stem-root. An overview of all 21 variables, their types, and possible values is provided in Table 2.

Notably, several gill-related and stem-related variables include "f" as a possible value, meaning "none". This raised a question about whether "f" represents a valid value or missing data. Analysis revealed that if any gill-related variable has a value of "f", all other gill-related variables for that mushroom also have a value of "f". A similar pattern was found for stem-related variables. Overall, 3,414 mushrooms lacked gill-related information, and 915 lacked stem-related information. These 915 mushrooms also had stem-height and stem-width equal to 0, confirming the consistency of this pattern. Gaining some domain knowledge revealed that some mushrooms

Variable	Missing Values	
	Count	Percentage
cap-surface	14 120	23.12
gill-attachment	9855	16.14
gill-spacing	25 062	41.04
stem-root	51 536	84.39
stem-surface	38 122	62.42
veil-type	57 746	94.56
veil-color	53 510	87.62
ring-type	2471	4.05
spore-print-color	54 597	89.40

Table 2: Mushroom Dataset Variables

Variable	Type	Possible Values
class	categorical	e (edible), p (poisonous/of unknown edibility)
cap-diameter	numerical	float number in cm
cap-shape	categorical	b (bell), c (conical), x (convex), f (flat), s (sunken), p (spherical), o (others)
cap-surface	categorical	i (fibrous), g (grooves), y (scaly), s (smooth), h (shiny), l (leathery), k (silky), t (sticky), w (wrinkled), e (fleshy), d (not specified by the author)
cap-color	categorical	n (brown), b (buff), g (gray), r (green), p (pink), u (purple), e (red), w (white), y (yellow), l (blue), o (orange), k (black)
does-bruise-or-bleed	categorical	t (bruises or bleeding), f (no)
gill-attachment	categorical	a (adnate), x (adnexed), d (decurrent), e (free), s (sinuate), p (pores), f (none)
gill-spacing	categorical	c (close), d (distant), f (none)
gill-color	categorical	n (brown), b (buff), g (gray), r (green), p (pink), u (purple), e (red), w (white), y (yellow), o (orange), k (black), f (none)
stem-height	numerical	float number in cm
stem-width	numerical	float number in mm
stem-root	categorical	b (bulbous), s (swollen), c (club), r (rooted), f (not specified by the author)
stem-surface	categorical	i (fibrous), y (scaly), s (smooth), h (shiny), k (silky), t (sticky), f (none), g (not specified by the author)
stem-color	categorical	n (brown), b (buff), g (gray), r (green), p (pink), u (purple), e (red), w (white), y (yellow), l (blue), o (orange), k (black), f (none)
veil-type	categorical	u (universal)
veil-color	categorical	n (brown), u (purple), e (red), w (white), y (yellow), k (black)
has-ring	categorical	t (ring), f (none)
ring-type	categorical	e (evanescent), r (flaring), g (grooved), l (large), p (pendant), z (zone), m (movable), f (none)
spore-print-color	categorical	n (brown), g (gray), r (green), p (pink), u (purple), w (white), k (black)
habitat	categorical	g (grasses), l (leaves), m (meadows), p (paths), h (heaths), u (urban), w (waste), d (woods)
season	categorical	s (spring), u (summer), a (autumn), w (winter)

naturally lack gills or a stem. This suggests that the values of "f" and 0 reflect genuine biological characteristics rather than missing data. Consequently, these values were treated as valid in subsequent analyses.

Lastly, the veil-type variable was examined for its potential utility in predictive modeling. This variable only takes the value "u (universal)", but it was missing for the majority (94.56%) of observations. Additionally, when veil-type was missing, both classes (edible and poisonous) were represented, indicating that the missingness did not provide meaningful predictive information for classification. As a result, the veil-type variable was excluded from further analysis.

4 Model Training and Evaluation

The first approach to training the decision tree classifier utilized the full dataset without removing any columns or rows, taking advantage of the `DecisionTreeClassifier` class's capability to handle missing values. Hyperparameter tuning was conducted using nested cross-validation with five folds for both the outer and inner loops. The parameter grid for this process was as follows:

```
parameter_grid = {
    'min_samples_split': [2, 5, 10, 20],
    'max_depth': [5, 10, 15, 20, None],
    'n_features': ["log2", "sqrt", None],
    'criterion': ["gini", "scaled_entropy", "square_root"],
    'min_information_gain': [0.0, 0.01, 0.05, 0.1],
    'n_quantiles': [5, 10, 20],
    'isolate_one': [True, False]
}
```

To mitigate the computational burden of evaluating all 4,320 possible parameter combinations, 86 (2%) combinations were randomly selected for evaluation.

The results demonstrated strong predictive performance, with no evidence of overfitting. The model achieved a mean test error of 0.0098, corresponding to a mean accuracy of 99.02%. Mean precision, recall, and F1 score were also high, at 99.30%, 98.93%, and 99.11%, respectively. The best-performing hyperparameters, identified based on the lowest test error (0.0055), were the following:

- `min_samples_split=10,`
- `max_depth=20,`
- `n_features=None,`
- `criterion='square_root',`
- `min_information_gain=0.0,`
- `n_quantiles=10,`
- `isolate_one=False`

A model trained with these hyperparameters yielded strong results on the test fold, with an accuracy of 99.45%, precision of 99.58%, recall of 99.42%, and an F1 score of 99.50%.

To investigate the influence of missing data on these results, a second approach was explored. Features with more than 40% missing values were removed, and rows containing any remaining missing values were excluded. This preprocessing step reduced the dataset to 36,948 observations and 14 features, with a class distribution of 16,944 edible mushrooms (45.86%) and 20,004 poisonous ones (54.14%).

The same nested cross-validation procedure and parameter grid were applied to this cleaned dataset, evaluating 86 randomly selected parameter combinations. The model's performance improved across all metrics. The mean test error dropped significantly to 0.0030, corresponding to a mean accuracy of 99.70%. Mean precision, recall, and F1 score also improved, reaching 99.78%, 99.67%, and 99.72%, respectively. The optimal hyperparameters, identified again based on the lowest test error (0.0022), were as follows:

- `min_samples_split=2,`
- `max_depth=None,`
- `n_features=None,`
- `criterion='scaled_entropy',`
- `min_information_gain=0.01,`
- `n_quantiles=20,`
- `isolate_one=True`

A decision tree trained with these hyperparameters achieved an accuracy of 99.78%, precision of 99.85%, recall of 99.75%, and an F1 score of 99.80% on the test fold.

Comparing these two approaches reveals the subtle yet significant impact of preprocessing on model performance. While the classifier's built-in capability to handle missing values was effective, the removal of features with a high proportion of missing values, combined with the elimination of incomplete rows, led to modest but meaningful improvements in predictive accuracy and reliability.

5   Conclusions

This project successfully implemented a custom decision tree classifier for mushroom classification, demonstrating high predictive performance. Nested cross-validation ensured robust hyperparameter tuning and evaluation, while preprocessing techniques enhanced model accuracy and reliability.

Future work could address some limitations by exploring more realistic datasets, employing advanced methods for handling missing values, and incorporating metrics for interpretability and computational efficiency. In addition, integrating ensemble methods, such as random forests or gradient boosting, could provide further insights into the potential of tree-based models for mushroom classification and similar tasks.